
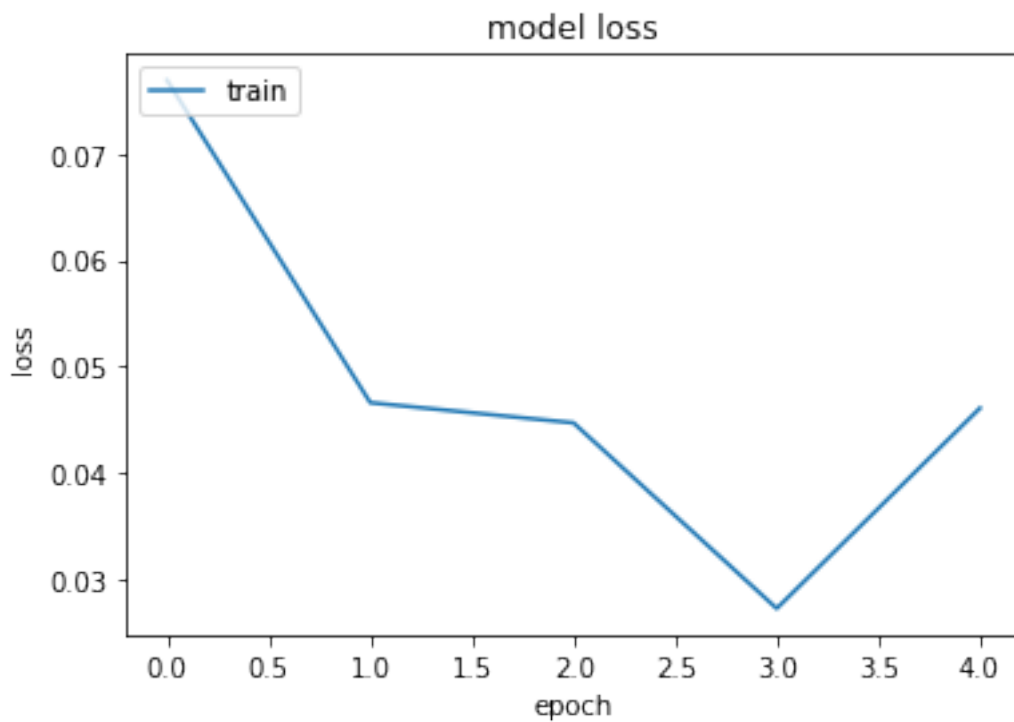


BERT:

17/17 [=====] - 4s 229ms/step - loss: 1.5301 - accuracy: 0.7897
[1.5301103591918945, 0.7896749377250671]

The loss of the model on the test set : 1.5301103591918945

The accuracy of the model on the test set : 78.96749377250671%



```
[45] new_input = ["i have a big butt"]  
  
     new_val = bert_encode(new_input, tokenizer, max_len=max_len)  
  
     print(model.predict(new_val))  
  
[[1.8674208e-06]]
```

```
[47] new_input = ["I fantasise about big butts at work"]  
  
     new_val = bert_encode(new_input, tokenizer, max_len=max_len)  
  
     print(model.predict(new_val))  
  
[[0.99968445]]
```

```
▶ new_input = ["I want to see your big boobies"]  
  
     new_val = bert_encode(new_input, tokenizer, max_len=max_len)  
  
     print(model.predict(new_val))  
  
☞ [[0.9999051]]
```

```
[50] new_input = ["boobies is an offensive word"]  
  
     new_val = bert_encode(new_input, tokenizer, max_len=max_len)  
  
     print(model.predict(new_val))  
  
[[2.1722178e-06]]
```

```
[42] new_input = ["lick my hairy balls"]  
  
     new_val = bert_encode(new_input, tokenizer, max_len=max_len)  
  
     print(model.predict(new_val))  
  
[[0.99990785]]
```

```
[37] new_input = ["i wanna eat your ass"]

new_val = bert_encode(new_input, tokenizer, max_len=max_len)

print(model.predict(new_val))

[[0.9999126]]
```

```
▶ new_input = ["bite me in the ass"]

new_val = bert_encode(new_input, tokenizer, max_len=max_len)

print(model.predict(new_val))

☞ [[0.9996673]]
```

```
▶ new_input = ["i am an ass"]

new_val = bert_encode(new_input, tokenizer, max_len=max_len)

print(model.predict(new_val))

☞ [[2.245538e-06]]
```

```
▶ new_input = ["i am a hairy person and i like playing with tennis balls"]

new_val = bert_encode(new_input, tokenizer, max_len=max_len)

print(model.predict(new_val))

☞ [[8.212096e-05]]
```

```
[39] new_input = ["a donkey can also be called an ass"]

new_val = bert_encode(new_input, tokenizer, max_len=max_len)

print(model.predict(new_val))

[[0.00032227]]
```

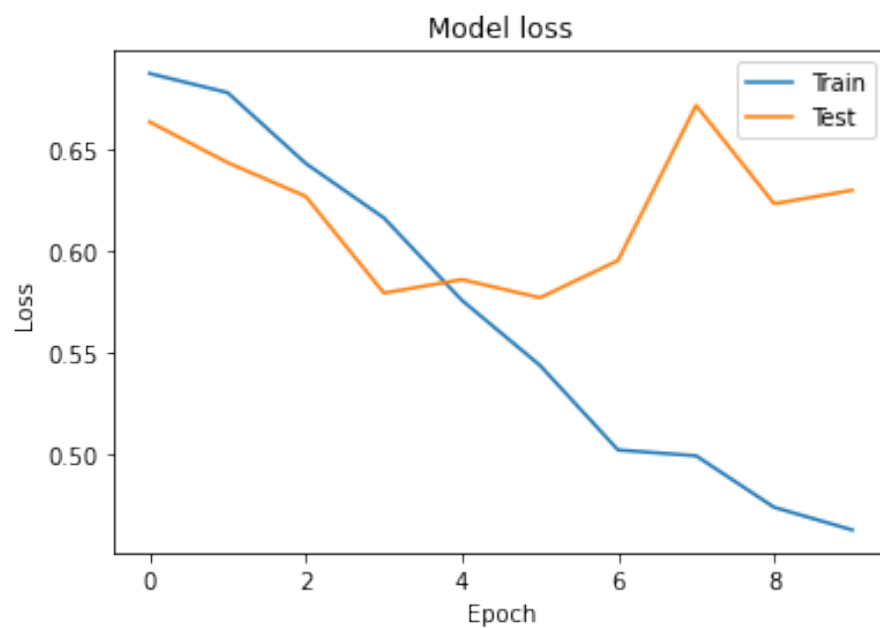
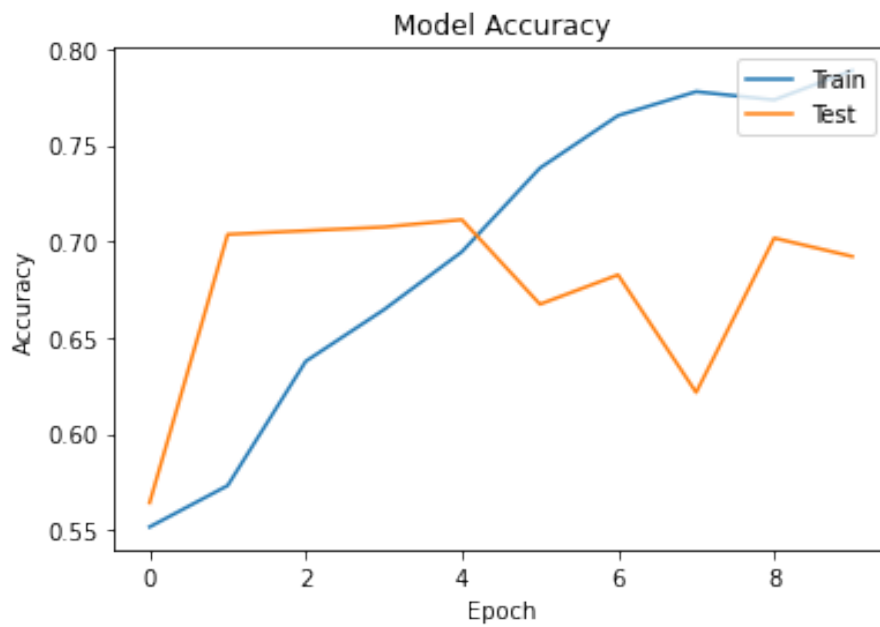
CNN-BiLSTM:

```
[89] test_performance=model.evaluate(x_test,y_test)
print("The loss of the model on the test set : {}".format(test_performance[0]))
print("The accuracy of the model on the test set : {}".format(test_performance[1]*100))
```

```
17/17 [=====] - 0s 14ms/step - loss: 0.6295 - accuracy: 0.6922
The loss of the model on the test set : 0.6294759511947632
The accuracy of the model on the test set : 69.21606063842773%
```

☞ Model: "sequential_11"

Layer (type)	Output Shape	Param #
=====		
embedding_7 (Embedding)	(None, 300, 300)	60000000
conv1d_11 (Conv1D)	(None, 298, 32)	28832
max_pooling1d_11 (MaxPooling)	(None, 149, 32)	0
dropout_22 (Dropout)	(None, 149, 32)	0
bidirectional_11 (Bidirectio	(None, 149, 256)	164864
flatten_11 (Flatten)	(None, 38144)	0
dense_22 (Dense)	(None, 30)	1144350
dropout_23 (Dropout)	(None, 30)	0
dense_23 (Dense)	(None, 2)	62
=====		
Total params: 7,338,108		
Trainable params: 1,338,108		
Non-trainable params: 6,000,000		



ML CLASSIFIERS:

Linear SVC:

0.9784585926280517

0.7571701720841301

Gaussian NB:

0.8879846816658689

0.6347992351816444

Logistic Regression:

0.8736237434179033

0.7571701720841301

Random Forest:

0.9784585926280517

0.7112810707456979